

Bottom- k Sampling, Frequency, and Set Similarity

Advanced Algorithms 2014, Assignment 2

Mikkel Thorup

This is the second mandatory assignment for the course Advanced Algorithms 2014. The Assignment should be solved in groups of 2 or 3 students. Permission for individual assignments is only granted in special cases. The assignment is due Sunday May 25th at 22:00. To pass the assignment the group must have completed most of the questions satisfactory. To be allowed to resubmit the group must have made a reasonable attempt at answering most of the questions. The resubmission is due Monday June 2nd at 22:00. Completed assignments must be uploaded to Absalon (this page). All descriptions and arguments should be kept concise while still containing relevant points.

1 Hash functions for sampling

The simplest use of hashing for sampling is to pick all keys that hash below a certain threshold (c.f. [1, §3.1]). In the context of sampling, it is often convenient to consider hash functions $h : U \rightarrow [0, 1)$ mapping the key universe U to the unit interval $[0, 1)$, that is, for any $x \in U$, $0 \leq h(x) < 1$. In practice we may have a strongly universal hash function $h_m : U \rightarrow [m]$ for some large m , and then we define $h(x) = h_m(x)/m$.

Exercise 1 (a) If $p \geq 100/m$, how far can $\Pr[h_m(x)/m < p]$ be from p ? Could we have equality?

(b) If $A \subseteq U$ and $m \geq 100|A|^2$, bound the probability that there are two keys $x, y \in A$ which get the same hash value $h_m(x)/m = h_m(y)/m$.

Below, for simplicity, we assume that we have a strongly universal hash function $h : U \rightarrow [0, 1)$. In particular we assume that h is collision free and that $\Pr[h(x) < p] = p$ for any given $p \in [0, 1)$. Moreover, we assume independent hash values $h(x)$ and $h(y)$ for any pair distinct keys $x, y \in U$. Define the sample of a key set $A \subseteq U$ as

$$S_{h,p}(A) = \{x \in A \mid h(x) < p\}.$$

For a given (non-random) p , we refer to this as *threshold sampling*. As described in [1, §3.1], we get

Lemma 1 For a given $p \in [0, 1)$ and set $A \subseteq U$, $n = |A|$, let $X = |S_{h,p}(A)|$. Then $\mu = \mathbb{E}[X] = pn$ and $\sigma^2 = \text{Var}[X] = (1 - p)\mu \leq \mu$. In particular, for any $r > 0$,

$$\Pr[|X - \mu| \geq r\sqrt{\mu}] \leq 1/r^2. \quad (1)$$

2 Bottom- k sampling

One issue with the above threshold sampling is that the number of samples is a variable depending on the size of the set A that we sample from. In many applications, we want to specify a hard limit k on the number of samples. A simple solution is to store a so-called *bottom- k sample*:

$$S_h^k(A) = \{\text{the } k \text{ keys } x \in A \text{ with the smallest hash values}\}$$

Here we assume that A has at least k keys and that there are no collisions between hash values from A .

If h was a truly random hash function, then $S_h^k(A)$ would be a uniformly random subset of A of size k . The reason is that the hash values would be distributed randomly between the keys in A , so any size- k subset would have exactly the same probability of getting the k smallest hash values. In particular each $x \in A$ would have exactly the same probability $p = k/n$ of belonging to $S_h^k(A)$.

2.1 Frequency estimation

Often the point of sampling is that we later want to estimate the frequency of a subset $C \subseteq A$ as the frequency of C among the k samples. That is, we estimate the frequency $f = |C|/|A|$ as $|C \cap S_h^k(A)|/k$.

Exercise 2 Assuming that $S_h^k(A)$ is a uniformly random size- k subset of A , prove that $E[|C \cap S_h^k(A)|/k] = |C|/|A|$.

Exercise 3 A common context in which bottom- k samples is applied is that the keys from A arrive online as a stream x_1, \dots, x_n .

- (a) What kind of data structure would you use to maintain the bottom- k sample as the keys arrive, that is, when you have received keys x_1, \dots, x_i , you should have their sample $S_h^k(\{x_1, \dots, x_i\})$?
- (b) How long time would it take you to process the next key x_{i+1} .

2.2 Similarity estimation

As discussed in [1, Section 3.1], one of the important points in sampling is that we want to compare sets A and B via their samples. We will use bottom- k samples to estimate their so-called Jaccard similarity $|A \cap B|/|A \cup B|$. This is the frequency of the intersection $A \cap B$ inside the union $A \cup B$.

Exercise 4 Assume that we have the bottom- k samples $S_h^k(A)$ and $S_h^k(B)$.

- (a) Prove that $S_h^k(A \cup B) = S_h^k(S_h^k(A) \cup S_h^k(B))$.
- (b) Prove that $A \cap B \cap S_h^k(A \cup B) = S_h^k(A) \cap S_h^k(B) \cap S_h^k(A \cup B)$.
- (c) Assuming that h is truly random and collision free, it now follows from Exercise 2 that

$$|S_h^k(A) \cap S_h^k(B) \cap S_h^k(S_h^k(A) \cup S_h^k(B))|/k$$

is an unbiased estimator of the Jaccard similarity $|A \cap B|/|A \cup B|$. How long time would it take you to compute the above estimate from $S_h^k(A)$ and $S_h^k(B)$. You may assume that $S_h^k(A)$ and $S_h^k(B)$ are sorted according to hash value.

3 Bottom- k sampling with strong universality

We are now going to bound the error probability of bottom- k estimates when based on a strongly universal hash function $h : U \rightarrow [0, 1]$. We are given the bottom- k sample $S = S_h^k(A)$ of a set A . We wish to use the sample S to estimate the frequency $f = |C|/|A|$ of a given subset $C \subseteq A$ as $|C \cap S|/k$. Using that $h : U \rightarrow [0, 1]$ is strongly universal, we will prove for any $r \leq \bar{r} = \sqrt{k}/3$ that

$$\Pr \left[|C \cap S|/k > f + 3r\sqrt{f/k} \right] \leq 2/r^2. \quad (2)$$

Note how increasing the number k of samples decreases the error $3r\sqrt{f/k}$.

There is a symmetric bound for under estimates

$$\Pr \left[|C \cap S|/k < f - 3r\sqrt{f/k} \right] \leq 2/r^2,$$

but it will not be proved during this assignment.

3.1 A union upper bound

For positive parameters $a < 1$ and b to be chosen later, we will bound the probability of the overestimate

$$|C \cap S| > \frac{1+b}{1-a} fk. \quad (3)$$

Define the threshold probability

$$p = \frac{k}{n(1-a)}.$$

Note that p is defined deterministically, independent of any samples. You will prove that the overestimate (3) implies one of the following two threshold sampling events:

- (I) The number of elements from A that hash below p is less than k .
- (II) The number of elements from C that hash below p is more than $(1+b)p|C|$.

Exercise 5 Prove that (I) and (II) are both false, then so is (3). As a first step, note that when (I) is false, all elements from the bottom- k sample S must hash below p .

From Exercise 5, we conclude that

Proposition 2 The probability of the overestimate (3) is bounded by $P_I + P_{II}$ where P_I and P_{II} are the probabilities of the events (I) and (II), respectively.

Upper bound with 2-independence For any given $r \leq \sqrt{k}/3$, we will fix a and b to give a combined error probability of $2/r^2$. More precisely, we will fix $a = r/\sqrt{k}$ and $b = r/\sqrt{fk}$. This also fixes $p = k/(n(1-a))$. We note for later that $a \leq 1/3$ and $a \leq b$. This implies

$$(1+b)/(1-a) \leq (1+3b) = 1 + 3r/\sqrt{fk}. \quad (4)$$

In connection with (I) we study the number X_A of elements from A hashing below p . The mean is $\mu_A = pn = k/(1-a)$. Now (I) is true if and only if $X_A < k = \mu_A(1-a) = \mu_A(1-r/\sqrt{k})$.

Exercise 6 Use Lemma 1 to prove that

$$P_I = \Pr[X_A < k] < 1/r^2.$$

In connection with (II) we study the number X_C of elements from C hashing below p . The mean is $\mu_C = p|C| = pfn = fk/(1-a) > fk$. Now (II) is true if and only if $X_C > \mu_C(1+b)$ where $b = r/\sqrt{fk}$.

Exercise 7 Use Lemma 1 to prove that

$$P_{II} = \Pr[X_C > (1+b)\mu_C] < 1/r^2.$$

By Proposition 2 we conclude that the probability of (3) is bounded by $P_I + P_{II} \leq 2/r^2$.

Rewriting (3) with (4), we conclude that

$$\Pr\left[|Y \cap S| > fk + 3r\sqrt{fk}\right] \leq 2/r^2. \tag{5}$$

This bounds the probability of the positive error in (3).

References

- [1] M. Thorup. Fast hashing, 2014.