

Proactive Computer Security

Write What Where - Heap Buffer Overflows

Morten Shearman Kirkegaard
<moki@fabletech.com>

2014-06-02

A *Write-What-Where condition* is a vulnerability which gives the attacker the ability to **write** a **value** (what) of the attackers choosing at an **address** (where) of the attackers choosing.

```
struct information {  
    char real_name[64];  
    ...  
};  
  
struct login {  
    char login_name[64];  
    struct information *info;  
};  
  
struct login *l = ...;  
strcpy(l->login_name, XXX_login_name);  
strcpy(l->info->real_name, XXX_real_name);
```

If the attacker can write a value at a location of her choosing, she has won.

Question

What should we overwrite?

We can overwrite function pointers (e.g. GOT entries) or program specific global variables (e.g. `set is_authenticated = true`).

Dynamically allocated variables are stored on the heap.

If one heap variable is a buffer, and we can trick the program to write outside this buffer, we might be able to take control of the process.


This is a relatively new subject. The first exploit was published in 2000 by Solar Designer.

Since then, priorities of heap managers have gone from "speed" to "resilience to errors, speed."

The heap manager has some internal data structures to keep track of allocations.

These data structures are often stored on the heap itself, next to the allocated memory.

Internally `dlmalloc` uses a *malloc_chunk* structure for each allocation. It starts 8 bytes (on 32-bit systems) before the allocation returned to the program.




Since all allocations are 8-byte aligned, the lower 3 bits of the size are always 0.


Two of them are used for flags:

`PINUSE_BIT` is set if the previous chunk is in use.


`CINUSE_BIT` is set if the current chunk is in use.



If the chunk before the current one is allocated (the `PINUSE_BIT` is set), the first field of the chunk header belongs to the previous chunk, and must not be accessed.





If the chunk before the current one is free (the `PINUSE_BIT` is clear), the first field of the chunk header contains the size of the previous chunk.



If a chunk (of less than 256 bytes) is free, the first 8 bytes are used as pointers for a doubly-linked list of free chunks of that size.


If the free block is 256 bytes or more, it is kept in a tree of free chunks. We will ignore this.






If a chunk is freed, and both surrounding chunks are allocated, the freed chunk is simply put in the doubly-linked list of the given `smallbin`.

The `C` bit is cleared in the header, the `P` bit is cleared in the next header, and the `prev_size` is set to the size of the free chunk.



If a chunk is freed, and both surrounding chunks are allocated, the freed chunk is simply put in the doubly-linked list of the given smallbin.


The C bit is cleared in the header, the P bit is cleared in the next header, and the prev_size is set to the size of the free chunk.



If a chunk is freed, and the previous chunk is free, the two are coalesced into one larger chunk.

This is done by unlinking the previous chunk from its free-list, and linking the new — larger — chunk into another free-list.


The P bit is cleared in the next header, and the size and new prev_size are set to the total size of the free chunk.



If a chunk is freed, and the previous chunk is free, the two are coalesced into one larger chunk.

This is done by unlinking the previous chunk from its free-list, and linking the new — larger — chunk into another free-list.


The P bit is cleared in the next header, and the size and new prev_size are set to the total size of the free chunk.



If a chunk is freed, and the next chunk is free, the two are coalesced into one larger chunk.

This is done by unlinking the next chunk from its free-list, and linking the new — larger — chunk into another free-list.


The C bit is cleared in the header, and the size and prev_size are set to the total size of the free chunk.





If a chunk is freed, and the next chunk is free, the two are coalesced into one larger chunk.


This is done by unlinking the next chunk from its free-list, and linking the new — larger — chunk into another free-list.


The C bit is cleared in the header, and the size and prev_size are set to the total size of the free chunk.











Unlinking is the most exploited part of dlmalloc and most other heap managers.

```
mchunkptr F = P->fd;  
mchunkptr B = P->bk;  
...  
F->bk = B;  
B->fd = F;
```

We have two pointers, F and B. At address (F + 12) the heap manager writes B, and at address (B + 8) it writes F. If we control F and B, we control the target process.

```
0x0804c485 <free+991>:  mov     eax,DWORD PTR [ebp-0x14]
0x0804c488 <free+994>:  mov     eax,DWORD PTR [eax+0x8]
0x0804c48b <free+997>:  mov     DWORD PTR [ebp-0x5c],eax
0x0804c48e <free+1000>: mov     eax,DWORD PTR [ebp-0x14]
0x0804c491 <free+1003>: mov     eax,DWORD PTR [eax+0xc]
0x0804c494 <free+1006>: mov     DWORD PTR [ebp-0x60],eax
...
0x804c4cc <free+1062>:  mov     eax,DWORD PTR [ebp-0x5c]
0x804c4cf <free+1065>:  mov     edx,DWORD PTR [ebp-0x60]
0x804c4d2 <free+1068>:  mov     DWORD PTR [eax+0xc],edx
0x804c4d5 <free+1071>:  mov     eax,DWORD PTR [ebp-0x60]
0x804c4d8 <free+1074>:  mov     edx,DWORD PTR [ebp-0x5c]
0x804c4db <free+1077>:  mov     DWORD PTR [eax+0x8],edx
```

We have a function pointer at 0x11111111
and our shellcode at 0x22222222.

If we force `fd` to 0x11111105 and `bk` to 0x22222222 the `unlink`
will do the following:

```
Set *(fd + 0xC) to bk;  
*(0x11111105 + 0xC) to 0x22222222;  
*0x11111111 to 0x22222222.
```

```
Set *(bk + 0x8) to fd;  
*(0x22222222 + 0x8) to 0x11111111;  
*0x2222222A to 0x11111111.
```

We now have what's known as *an arbitrary 4-byte memory overwrite*.

The first overwrite will change the function pointer, so it points at our shellcode. This is good.

The second overwrite will change part of the shellcode to point to the function pointer. This is annoying.

If the shellcode is structured like this, the overwrite will not disturb the flow:

22222222	90	nop
22222223	90	nop
22222224	90	nop
22222225	90	nop
22222226	90	nop
22222227	90	nop
22222228	EB04	jmp short 0x2222222E
2222222A	41	inc ecx
2222222B	41	inc ecx
2222222C	41	inc ecx
2222222D	41	inc ecx
2222222E	90	nop ; code here

```
char *a, *b, *c;
```

```
a = malloc(20);
```

```
b = malloc(20);
```


```
c = malloc(20);
```

->

```
free(b);
```

```
strcpy(b, "AAAABBBB");
```

```
free(a);
```



```
char *a, *b, *c;
```

```
a = malloc(20);
```

```
b = malloc(20);
```


```
c = malloc(20);
```

```
free(b);
```

->

```
strcpy(b, "AAAABBBB");
```

```
free(a);
```




```
char *a, *b, *c;
```

```
a = malloc(20);
```

```
b = malloc(20);
```


```
c = malloc(20);
```

```
free(b);
```

```
strcpy(b, "AAAABBBBB");
```

->

```
free(a);
```



Program received signal SIGSEGV, Segmentation fault.

0x0804c465 in free ()

(gdb) x/i \$eip

0x804c465 <free+1135>: mov WORD PTR [eax+0xc],edx

(gdb) i r eax edx

eax	0x41414141	1094795585
-----	------------	------------

edx	0x42424242	1111638594
-----	------------	------------

```
char *a, *b;
```

```
a = malloc(20);
```

```
b = malloc(20);
```


->

```
strcpy(a, "aaaabbbbccccdddd"
```

```
"eeee\xff\xff\xff\xff"
```

```
"AAAABBBB");
```

```
free(a);
```



```
char *a, *b;
```

```
a = malloc(20);
```

```
b = malloc(20);
```

```
strcpy(a, "aaaabbbbccccdddd"
```

```
"eeee\xff\xff\xff\xf1"
```

```
"AAAABBBB");
```

```
->
```

```
free(a);
```

size				1	1
61	61	61	61		
62	62	62	62		
63	63	63	63		
64	64		
FF	FF	FF		0	1
41	41	41	41		
42	42	42	42		
b					

Program received signal SIGSEGV, Segmentation fault.

0x0804c465 in free ()

(gdb) x/i \$eip

0x804c465 <free+1135>: mov DWORD PTR [eax+0xc],edx

(gdb) i r eax edx

eax	0x41414141	1094795585
-----	------------	------------

edx	0x42424242	1111638594
-----	------------	------------

```
char *a_in, *a_out;
```

```
char *b_in, *b_out;
```

```
a_in = "aaaabbbbccccddddeeee";
```

```
b_in = "AAAABBBB";
```

```
a_out = malloc(strlen(a_in));
```


```
b_out = malloc(strlen(b_in));
```

->


```
strcpy(a_out, a_in);
```

```
strcpy(b_out, b_in);
```

```
free(a_out);
```



```
char *a_in, *a_out;  
char *b_in, *b_out;  
  
a_in = "aaaabbbbccccddddeeee";  
b_in = "AAAABBBB";  
  
a_out = malloc(strlen(a_in));  
b_out = malloc(strlen(b_in));  
  
strcpy(a_out, a_in);  
-> strcpy(b_out, b_in);  
  
free(a_out);
```



```
char *a_in, *a_out;  
char *b_in, *b_out;  
  
a_in = "aaaabbbbccccddddeeee";  
b_in = "AAAABBBB";  
  
a_out = malloc(strlen(a_in));  
b_out = malloc(strlen(b_in));  
  
strcpy(a_out, a_in);  
  
strcpy(b_out, b_in);  
  
->  
free(a_out);
```

size	1	1
61	61	61
62	62	62
63	63	63
64	64	..
size	0	0
41	41	41
42	42	42
b_out		

Program received signal SIGSEGV, Segmentation fault.

0x0804c465 in free ()

(gdb) x/i \$eip


0x804c465 <free+1135>: mov DWORD PTR [eax+0xc],edx

(gdb) i r eax edx

eax	0x41414141	1094795585
-----	------------	------------

edx	0x42424242	1111638594
-----	------------	------------

```
-> p = malloc(20);  
  
free(p);  
  
free(p);  
  
q = malloc(20);  
  
strcpy(q,  
      "AAAABBBB");
```



```
p = malloc(20);
```


```
free(p);
```

->

```
free(p);
```

```
q = malloc(20);
```

```
strcpy(q,  
      "AAAABBBB");
```



```
p = malloc(20);
```


```
free(p);
```

```
free(p);
```

->

```
q = malloc(20);
```

```
strcpy(q,  
      "AAAABBBB");
```



```
p = malloc(20);
```


```
free(p);
```

```
free(p);
```

```
q = malloc(20);
```

->

```
strcpy(q,  
      "AAAABBBB");
```



```
p = malloc(20);
```


```
free(p);
```

```
free(p);
```

```
q = malloc(20);
```

```
strcpy(q,  
      "AAAABBBB");
```

->



Coffee Break

If you don't know exactly which address your shellcode will be placed at, you can make sure that it does not matter.

Make enough code, which does nothing. Place it in front of the shellcode.

When EIP is set to any address within this area, it will "slide" down to your shellcode.

A *NOP Slide* a/k/a *NOP Sled* a/k/a *NOP ...* can be as simple as:

90	nop
90	nop
90	nop
90	nop
XX XX XX XX XX	SHELLCODE

If it must allow overwrites (for unlink exploits), a series of jumps is useful:

```
EB 7E          jmp short +126
EB 7E          jmp short +126
EB 7E          jmp short +126
...
90             nop
90             nop
90             nop
...
XX XX XX XX XX SHELLCODE
```

If you have to guess a heap address of one buffer, and you must hit somewhere in a 1 kilobytes NOP Slide, your odds are small; About 1 in 4000000 on a 32 bit host.

If you have to guess a heap address of one buffer, and you must hit somewhere in a 1 kilobytes NOP Slide, your odds are small; About 1 in 4000000 on a 32 bit host.

If you have to guess the address of one of 1000 NOP Slides of each 1 megabyte, the odds are a lot better; About 1 in 4.

If the target process contains a scripting engine, like JavaScript or ActionScript, it can be used to do Heap Spraying.

```
var spraySlide = unescape("%u9090%u9090");  
while (spraySlide.length*2 < spraySlideSize) {  
    spraySlide += spraySlide;  
}
```

```
var memory = new Array();  
for (i=0; i<heapBlocks; i++) {  
    memory[i] = spraySlide + payload;  
}
```

NOP Slides and Heap Spraying are signs of unreliability.
Don't use them in the assignment.

```
char *buffer;  
void (**p_func_ptr)(void);
```


```
buffer = malloc(20);
```

```
p_func_ptr = malloc(4);  
*p_func_ptr = some_func;
```

->

```
strcpy(buffer,  
"AAAAAAAAA...");
```

```
(*p_func_ptr)();
```



```
char *buffer;
void (**p_func_ptr)(void);
```

```
buffer = malloc(20);
```

```
p_func_ptr = malloc(4);
*p_func_ptr = some_func;
```

```
strcpy(buffer,
  "AAAAAAAAA...");
```

->

```
(*p_func_ptr)();
```

size				1	1
41	41	41	41		
41	41	41	41		
41	41		
41	41	..		0	1
41	41	41	41		



```
Program received signal SIGSEGV, Segmentation fault.  
0x41414141 in ?? ()
```

```
class Demo {  
public:  
    virtual void method_a(void);  
    int attr_a;  
    int attr_b;  
};
```

```
char *buffer = new char[16];  
Demo *demo = new Demo;
```

->

```
strcpy(buffer,  
    "AAAAAAAAA...");  
  
demo->method_a();
```



```
class Demo {
public:
    virtual void method_a(void);
    int attr_a;
    int attr_b;
};
```

```
char *buffer = new char[16];
Demo *demo = new Demo;
```

```
strcpy(buffer,
    "AAAAAAAAA...");
```

->

```
demo->method_a();
```

size				1	1
41	41	41	41		
41	41	41	41		
41	41		
41	41	..	0	1	
41	41	41	41		
41	41	41	41		
41	41	41	41		

Program received signal SIGSEGV, Segmentation fault.

0x080486ce in main ()

(gdb) x/4i \$eip

0x80486ce <main+202>: mov edx,DWORD PTR [eax]

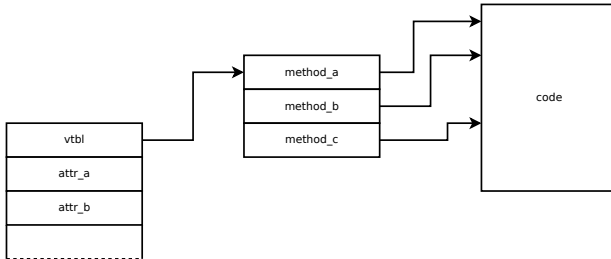
0x80486d0 <main+204>: mov eax,DWORD PTR [esp+0x18]

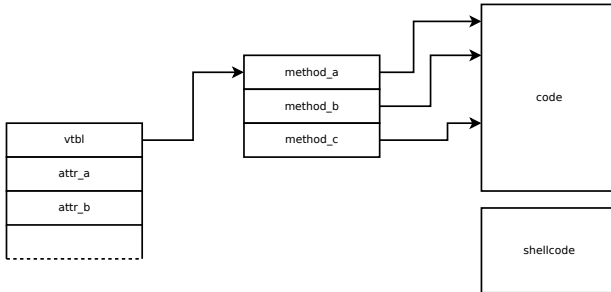
0x80486d4 <main+208>: mov DWORD PTR [esp],eax

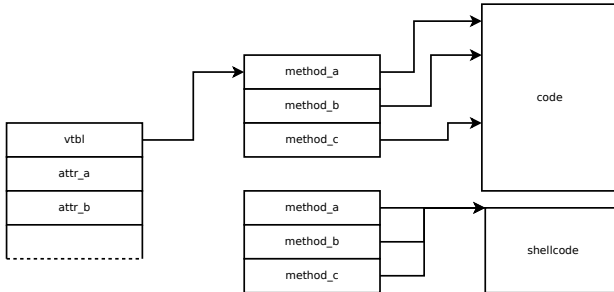
0x80486d7 <main+211>: call edx

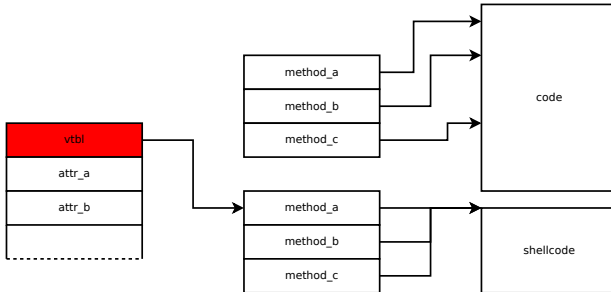
(gdb) i r eax

eax 0x41414141 1094795585





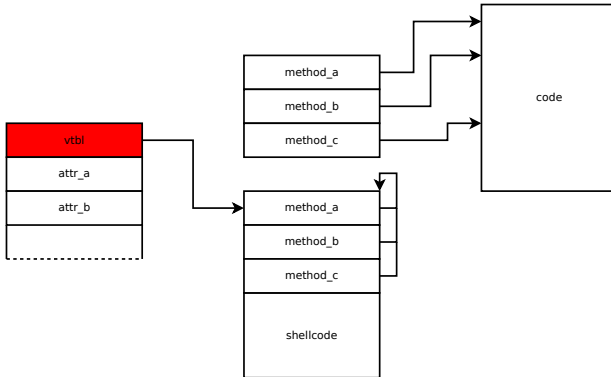




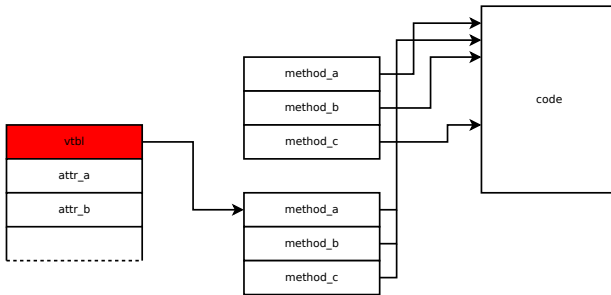
The address 0x0C0C0C0C is often used for spraying fake virtual method tables. You only need to spray around 200 MB, the alignment doesn't matter, and the address is valid code — almost NOPs.

0C0C0C0C	0C0C	or al,0xc
0C0C0C0E	0C0C	or al,0xc
0C0C0C10	0C0C	or al,0xc
0C0C0C12	0C0C	or al,0xc

NOP Sled and virtual method table in one.



With non-executable heap, the method pointers must point to the code segment, making exploitation more challenging.



In an exploit for a stack buffer overflow, the attacker can often set up a number of fake stack frames, chaining several *ROP gadgets*. This is often necessary because of ASLR.

In general, a heap buffer overflow "only" gives the attacker EIP control. It is not possible to chain ROP gadgets directly. The very first gadget *must* set up everything needed for chaining.

In this example, ESP points to the original stack, and EAX points to a heap buffer the attacker controls.

By copying EAX to ESP, the heap becomes the new stack, allowing the attacker to chain multiple ROP gadgets.

EIP → `mov esp, eax`
`ret`

ESP

→

...
7F 88 04 08
F8 CD FF FF
00 00 00 00
F4 4F FB F7

EAX →

29 AD 04 08
1E B1 04 08
41 41 41 41
C1 BB 04 08

In this example, ESP points to the original stack, and EAX points to a heap buffer the attacker controls.

By copying EAX to ESP, the heap becomes the new stack, allowing the attacker to chain multiple ROP gadgets.

```

    mov esp, eax
EIP ->  ret

```

ESP EAX ->

...
7F 88 04 08
F8 CD FF FF
00 00 00 00
F4 4F FB F7
29 AD 04 08
1E B1 04 08
41 41 41 41
C1 BB 04 08

Another instruction – XCHG – is very useful for making a heap buffer the new stack.

	...
	7F 88 04 08
	F8 CD FF FF
	00 00 00 00
ESP ->	F4 4F FB F7

EIP -> xchg esp, eax
ret


	29 AD 04 08
	1E B1 04 08
	41 41 41 41
EAX ->	C1 BB 04 08

Another instruction – XCHG – is very useful for making a heap buffer the new stack.


```
    xchg esp, eax  
EIP ->  ret
```

	...
EAX ->	7F 88 04 08
	F8 CD FF FF
	00 00 00 00
	F4 4F FB F7
	...
ESP ->	29 AD 04 08
	1E B1 04 08
	41 41 41 41
	C1 BB 04 08


Heap space has traditionally been allocated by moving the *program break*, using the BRK system call.



Heap space has traditionally been allocated by moving the *program break*, using the BRK system call.



With ASLR this gives some entropy, but not much. Only the base address of the heap is randomized.




OpenBSD's heap manager uses MMAP, rather than BRK, to allocate heap space.

This ensures that addresses are randomized, because of ASLR. It also prevents some overflow, use-after-free, and double-free bugs from being exploitable.

By using the MMAP system call to map new heap pages, a lot more entropy can be added to the heap layout. The price is performance.



By using the MMAP system call to map new heap pages, a lot more entropy can be added to the heap layout. The price is performance.

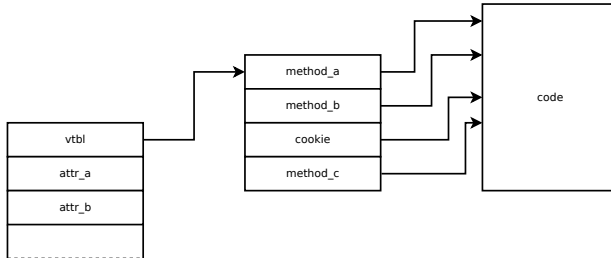


Today all heap managers have consistency checks. In dlmalloc it looks like this:

```
static void do_check_free_chunk(mstate m, mchunkptr p) {  
    ...  
    assert(!is_inuse(p));  
    assert(!next_pinuse(p));  
    ...  
    assert(next->prev_foot == sz);  
    assert(pinuse(p));  
    assert(next == m->top || is_inuse(next));  
    assert(p->fd->bk == p);  
    assert(p->bk->fd == p);  
}
```

Cisco IOS has a special process ("Check heaps") which verifies heap consistency.

According to Cisco's documentation it *[c]hecks the memory every minute. It forces a reload if it finds processor corruption.*



The Windows heap manager has an 8-bit cookie in each chunk header, and will terminate the process, if a cookie is invalid.

PHP has its own `malloc()` wrapper, which does not have *safe unlinking*. This disables the mitigations of the default heap manager — not a smart move.

The *Hardened PHP Project's Suhosin* patch adds this, as well as cookies to the header and footer of each heap chunk.

Remember to fill in the feedback form.

That's it. Have fun!